

Introduzione al Deep Learning

Giovanni Busonera

CRS4

17/05/2019

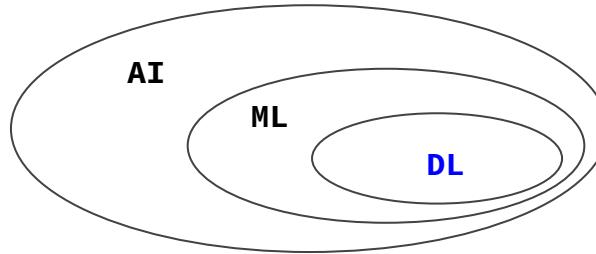


Deep Learning: Outline

1. *Definizione e contesto*
2. *Perceptron*
3. *Reti neurali*
4. *Introduzione a Keras*

Deep Learning:

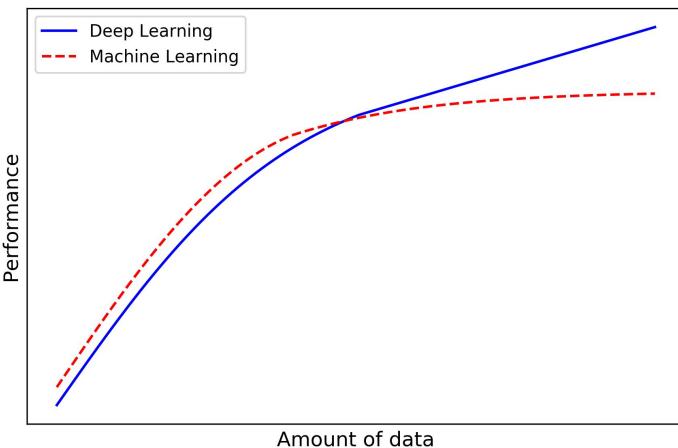
Il Deep Learning (DL) fa parte del Machine Learning (ML) e dell'Intelligenza Artificiale (AI). Si basa su diversi livelli di rappresentazione corrispondenti a gerarchie di concetti in cui quelli di alto livello sono definiti sulla base di quelli di basso livello mediante relazioni non lineari.



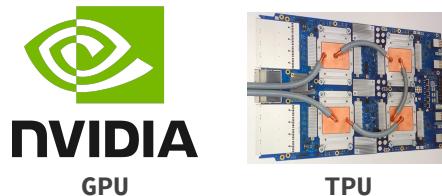
“Deep-learning methods are representation-learning methods with multiple levels of representation, obtained by composing simple but non-linear modules that each transform the representation at one level (starting with the raw input) into a representation at a higher, slightly more abstract level.”

Deep Learning: alcune differenze con il ML

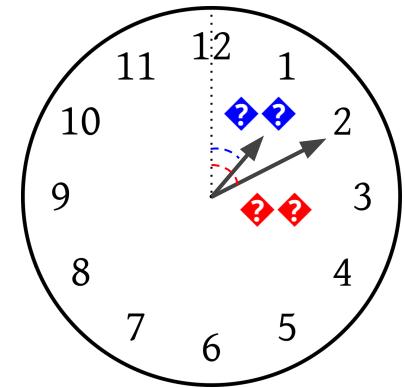
Prestazioni vs quantità di dati



Hardware

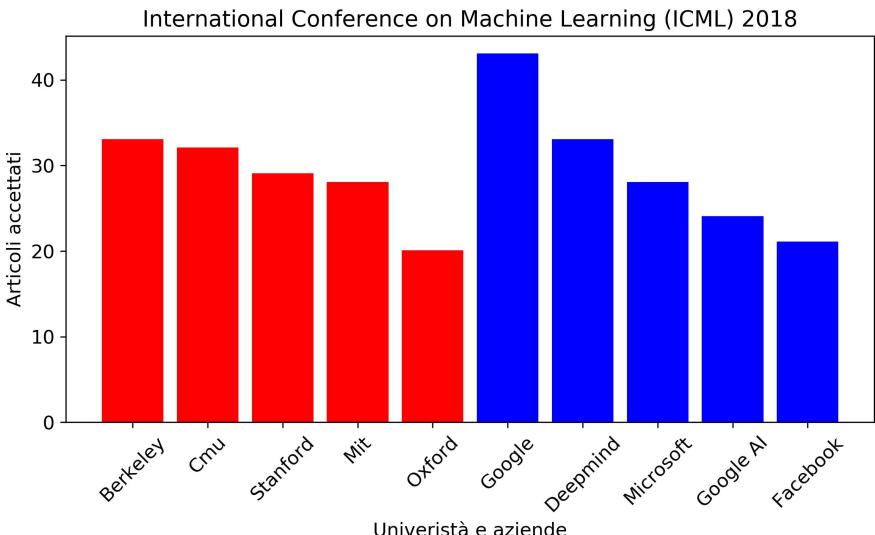
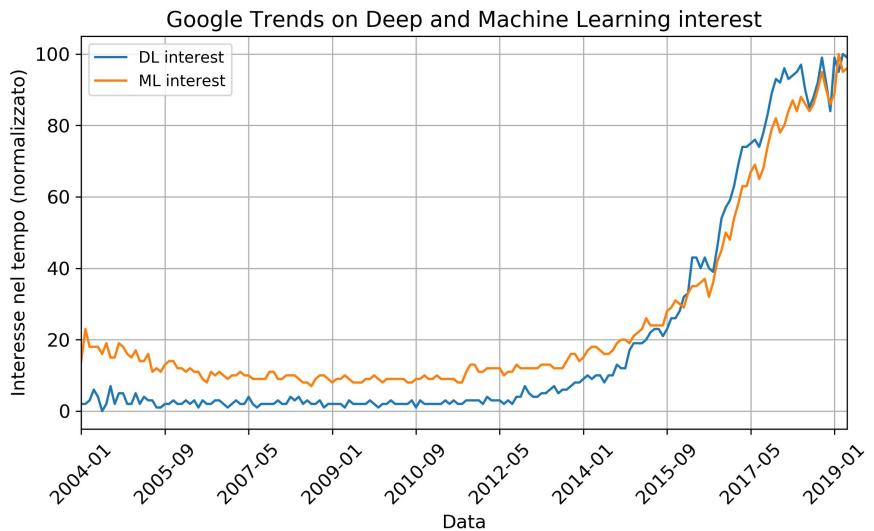


Feature extraction



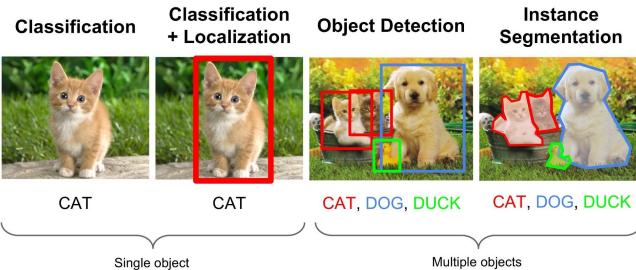
	Input	Target
ML	α, β	1:10AM
DL	whole image	1:10AM

Deep Learning: Trend



Applicazioni

Visual recognition



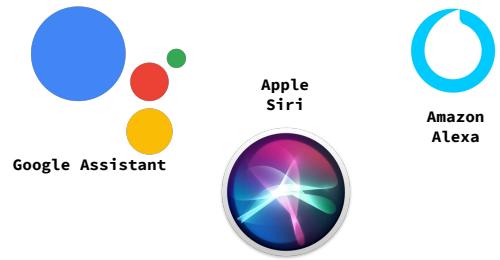
Self driving car



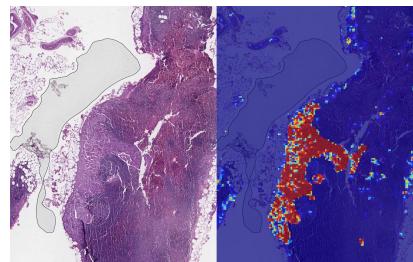
Search engine



Speech recognition



Healthcare



Multilingual translation

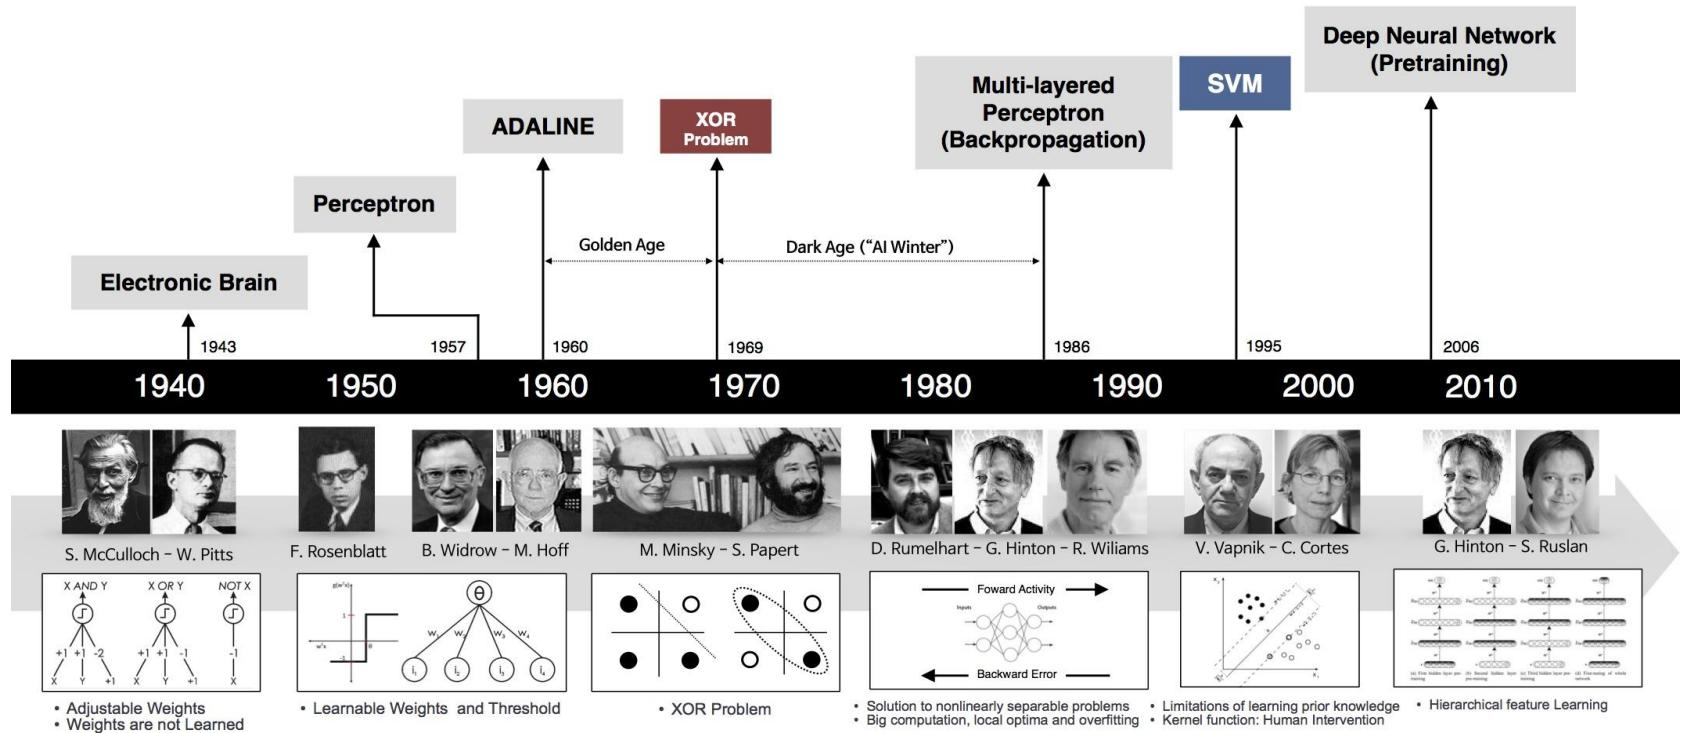


Google Translate



DeepL

Cenni storici: dal Perceptron al Deep Learning

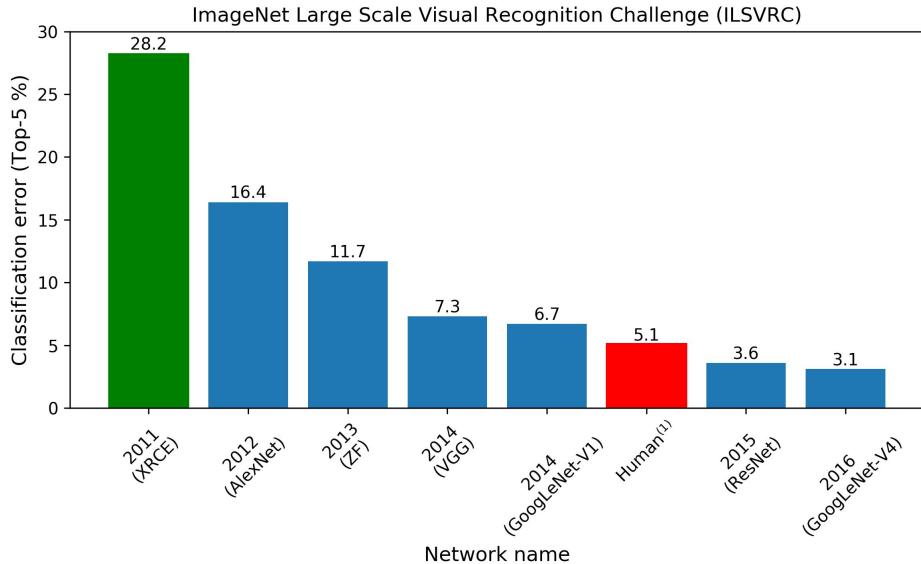


ImageNet e ILSVRC

ImageNet Challenge

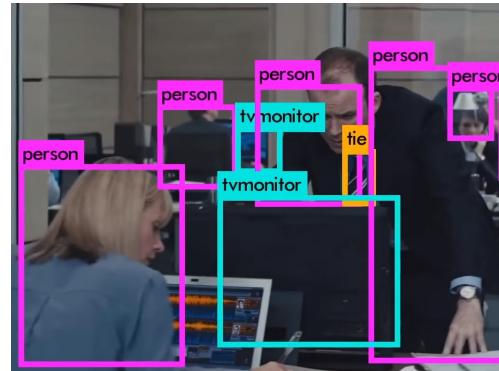


- 1,000 object classes (categories).
- Images:
 - 1.2 M train
 - 100k test.



(1) Russakovsky O. et Alter (2015). ImageNet Large Scale Visual Recognition Challenge. International Journal of Computer Vision, 115(3), 211-252

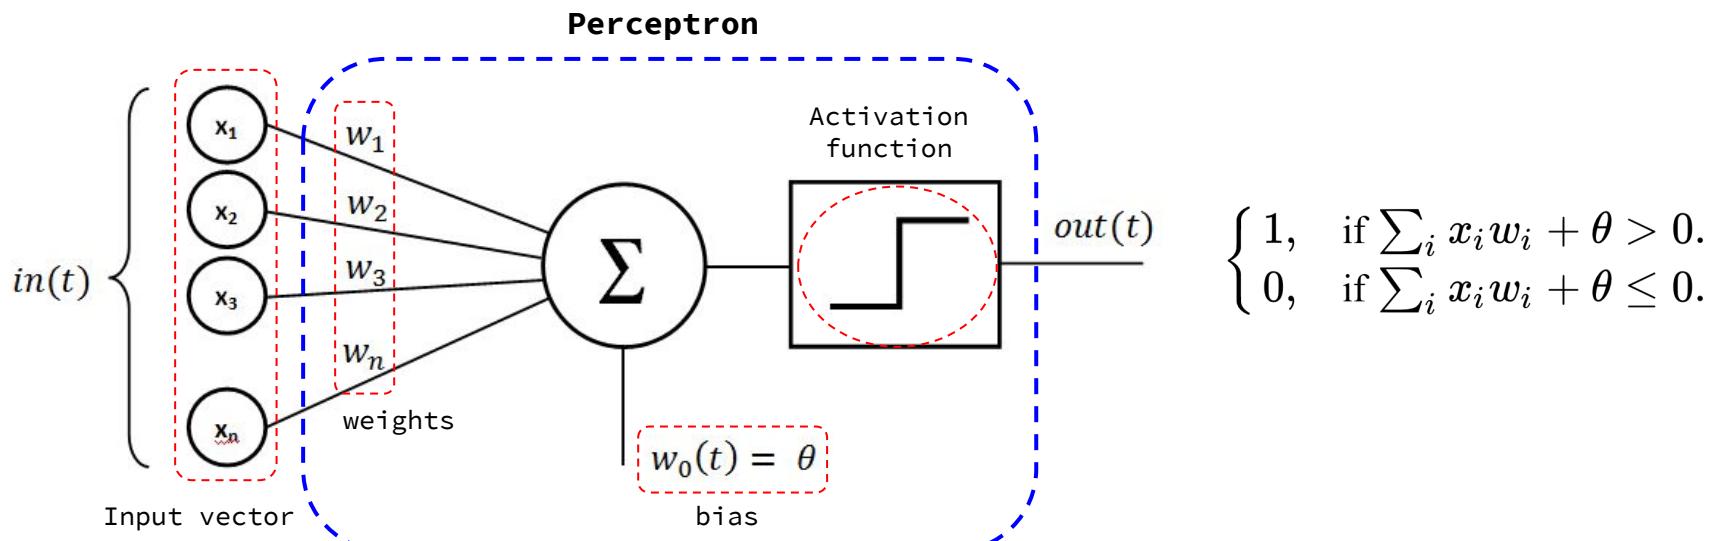
Real time object detection: Yolo Net v2



Perceptron

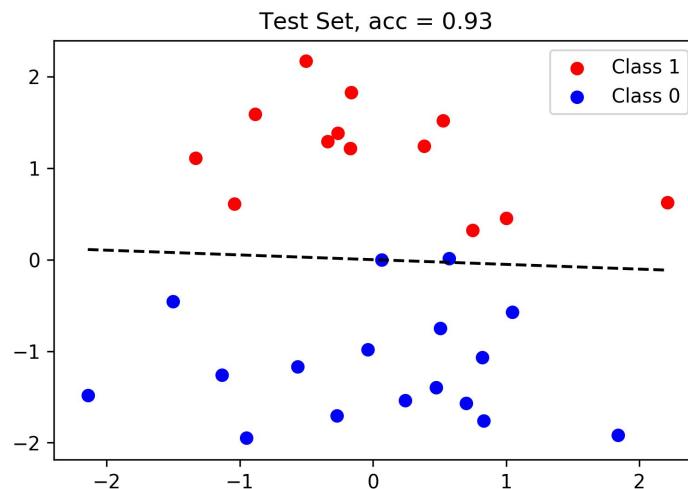
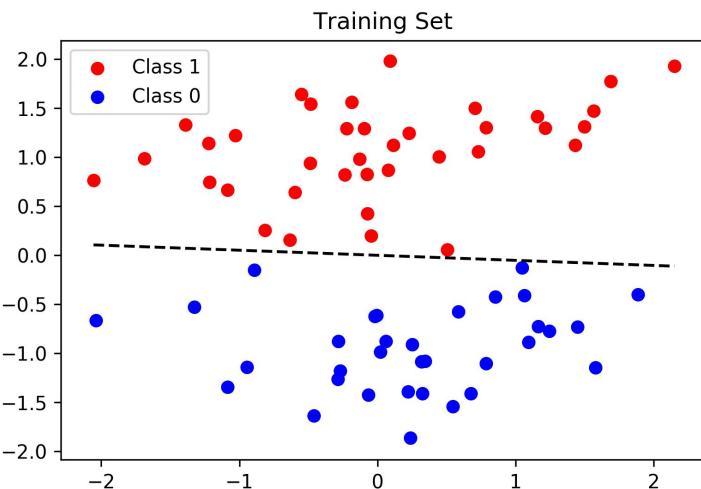
Origini delle ANN: il Perceptron

- È un classificatore lineare binario proposto da Frank Rosenblatt nel 1958 e ispirato dai precedenti lavori di Warren McCulloch e Walter Pitts.



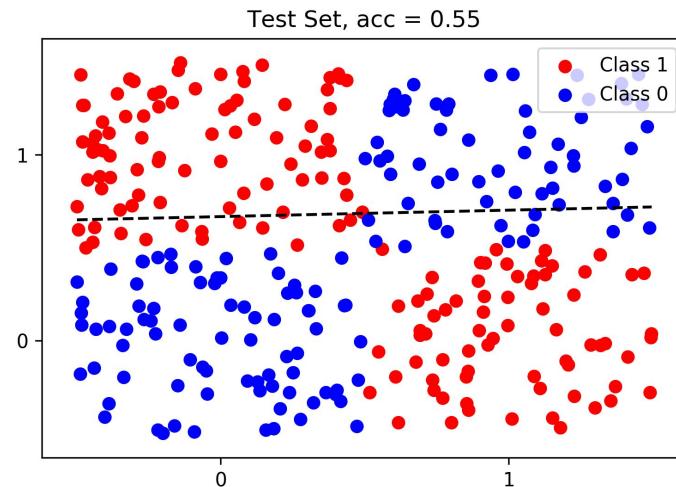
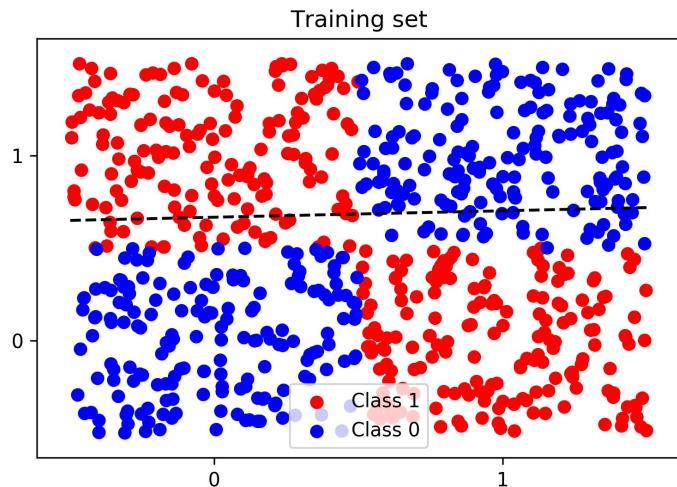
Origini delle ANN: il Perceptron

- L'algoritmo converge solo se l'insieme dei dati è linearmente separabile.



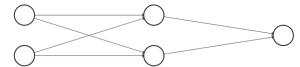
Origini delle ANN: il Perceptron e il problema dello XOR

- Per dati non separabili linearmente l'algoritmo da risultati insoddisfacenti

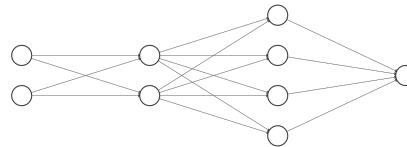


Origini delle ANN: La soluzione dello XOR → MLP

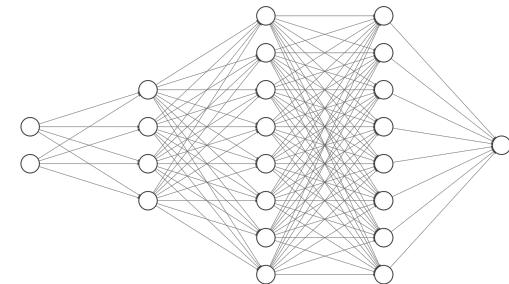
- Multilayer perceptron(MLP)



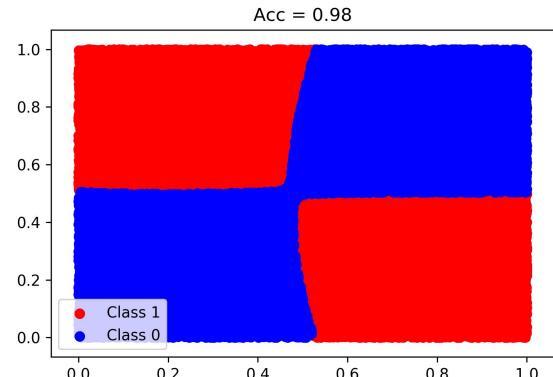
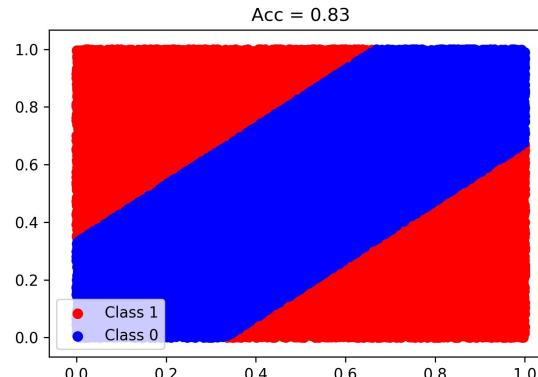
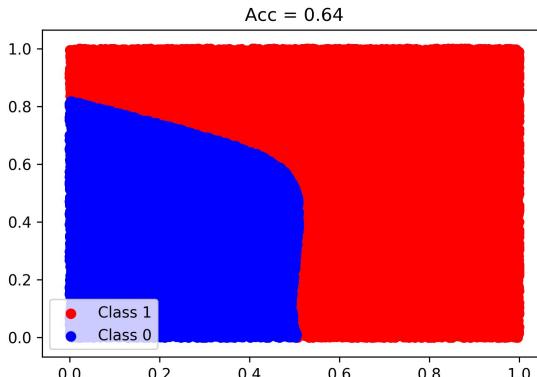
Input Layer $\in \mathbb{R}^2$ Hidden Layer $\in \mathbb{R}^2$ Output Layer $\in \mathbb{R}^1$



Input Layer $\in \mathbb{R}^2$ Hidden Layer $\in \mathbb{R}^2$ Hidden Layer $\in \mathbb{R}^4$ Output Layer $\in \mathbb{R}^1$



Input Layer $\in \mathbb{R}^2$ Hidden Layer $\in \mathbb{R}^4$ Hidden Layer $\in \mathbb{R}^8$ Hidden Layer $\in \mathbb{R}^8$ Output Layer $\in \mathbb{R}^1$



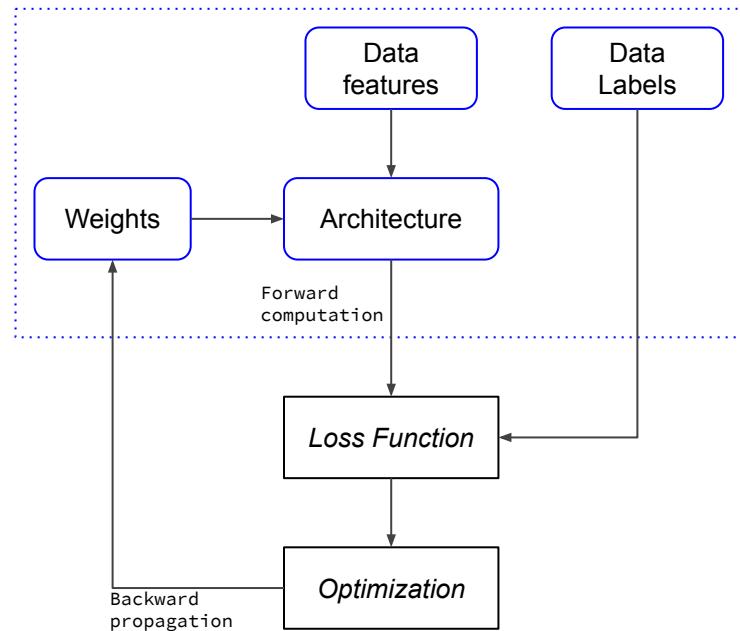
Origini delle ANN: L'apprendimento del MLP

- Aumentano i parametri e serve un metodo per determinarli in modo automatico (Gradient based optimization e backpropagation)
- La funzione di attivazione deve avere delle caratteristiche adatte (derivabile) per l'applicazione degli algoritmi di training

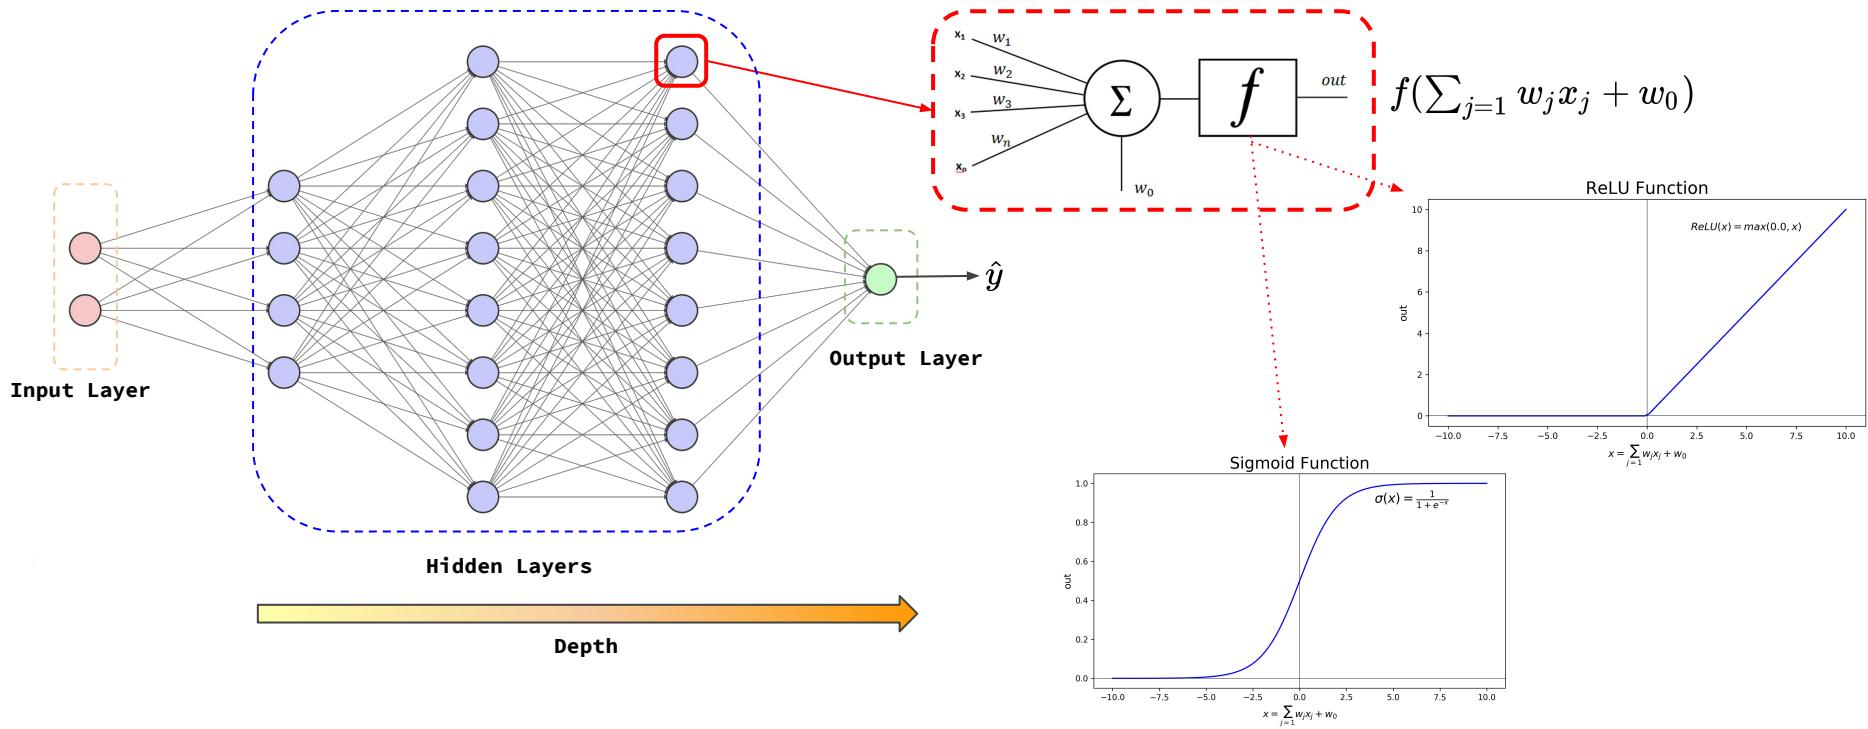
Artificial Neural Networks

Anatomia di una rete neurale

1. Architettura (forward computation)
2. Loss Function
3. Gradient based optimization (backward propagation)
4. Misura della prestazione (es: accuracy)



ANN: Architettura e Forward computation



Loss function

- Misura la “distanza” tra la predizione della rete neurale e l’output di riferimento dei samples. Ad esempio la Mean Squared Error (MSE) è definita come:

$$L(w) \equiv \frac{1}{n} \sum_i \|y_i - \hat{y}_i\|^2$$

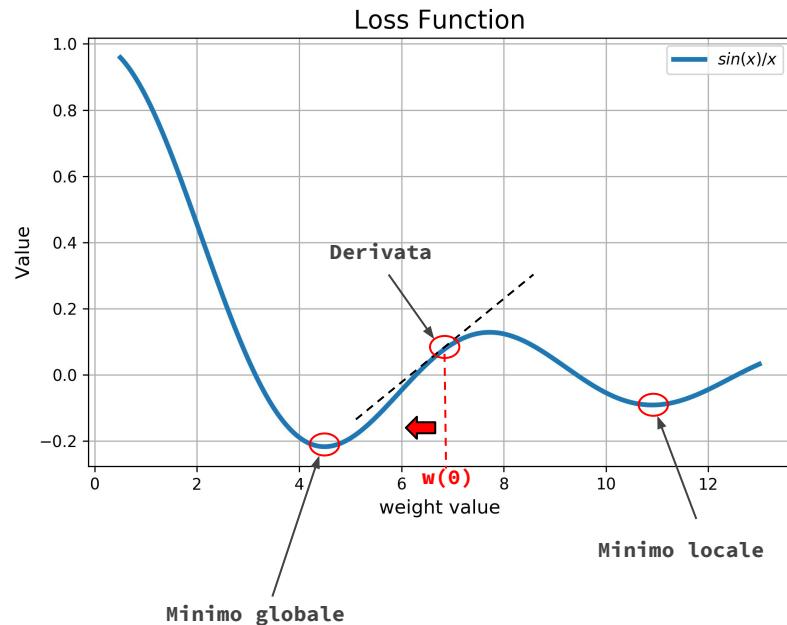
- Esistono varie tipologie di Loss function (MSE, Binary crossentropy, Categorical crossentropy, etc...)
- La scelta della Loss function dipende dall’applicazione

Optimization: Gradient Descent (GD)

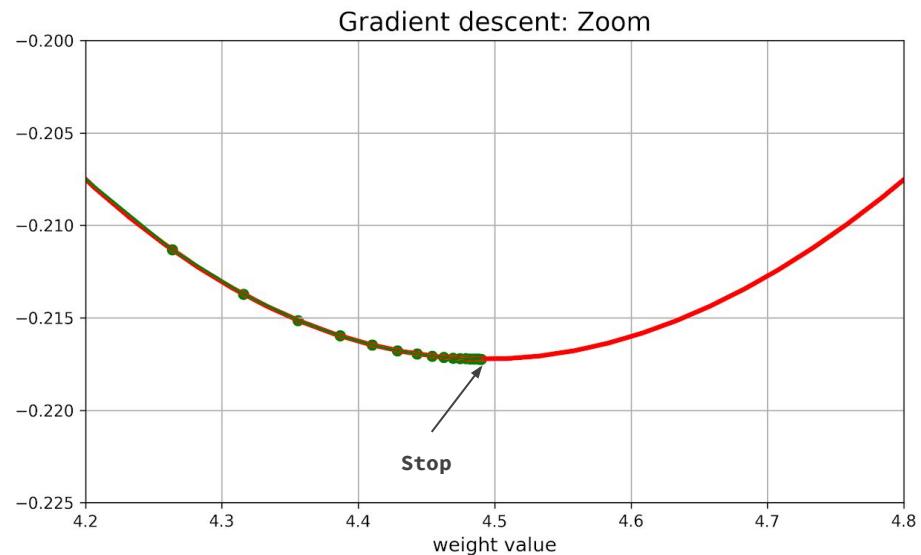
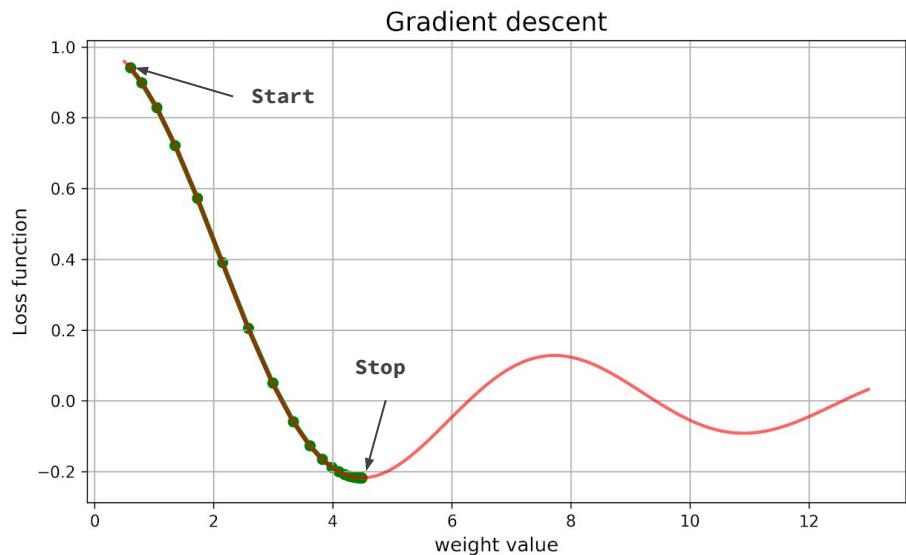
- Si sceglie un punto a caso $w(0)$
- Si calcola la derivata
- Ci si muove in direzione opposta alla derivata

$$w(i + 1) = w(i) - \eta \frac{dL}{dw}$$

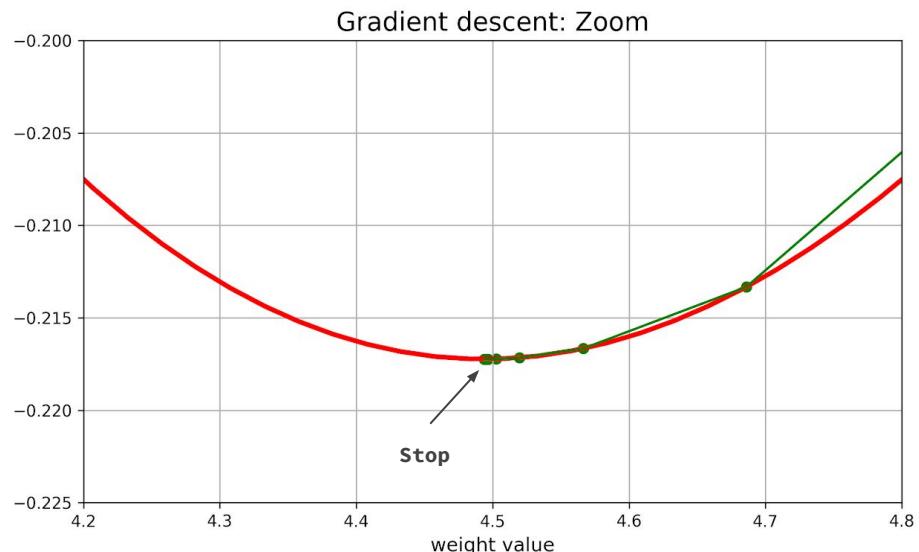
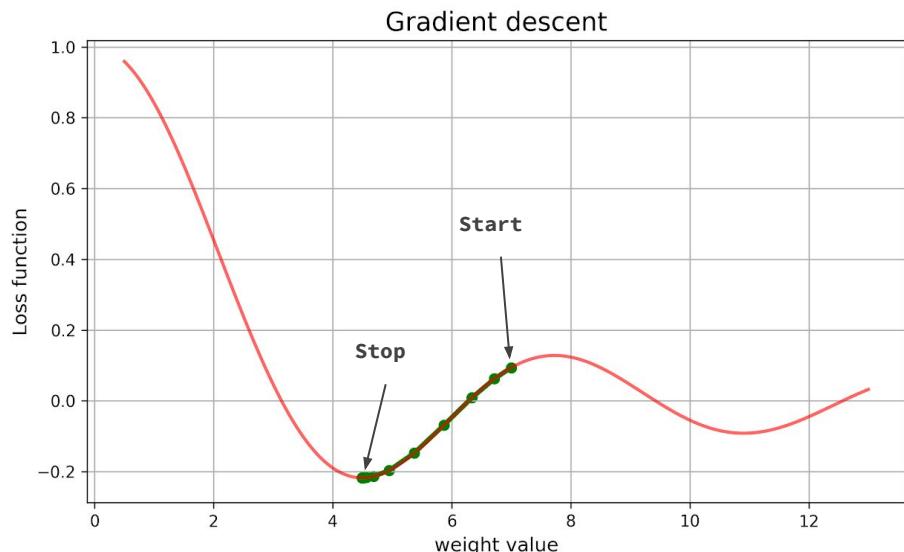
Learning rate



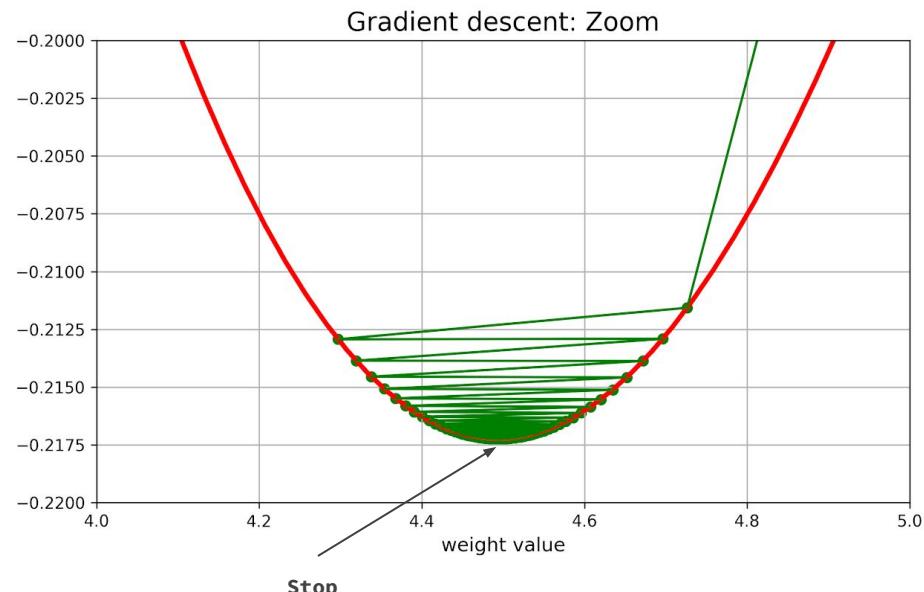
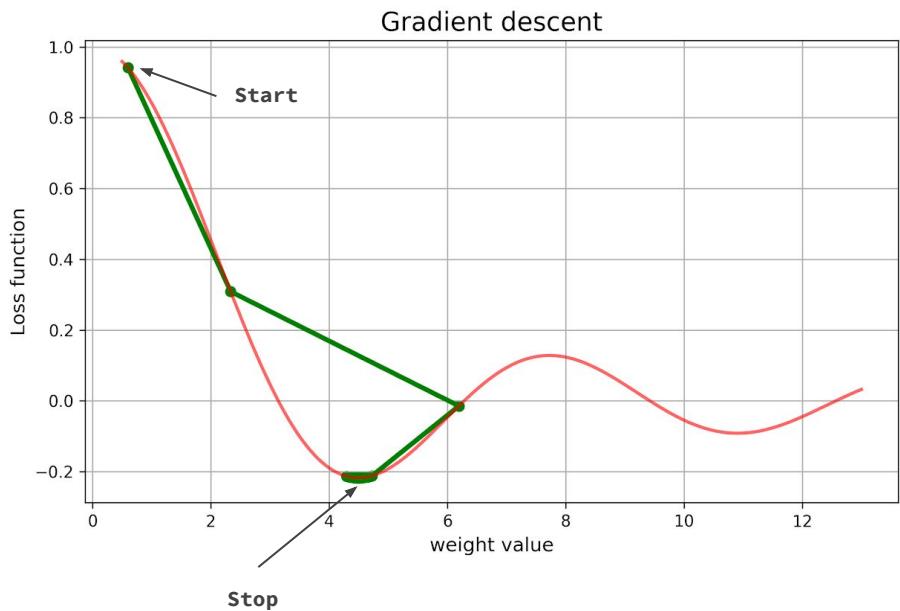
Optimization: Gradient descent



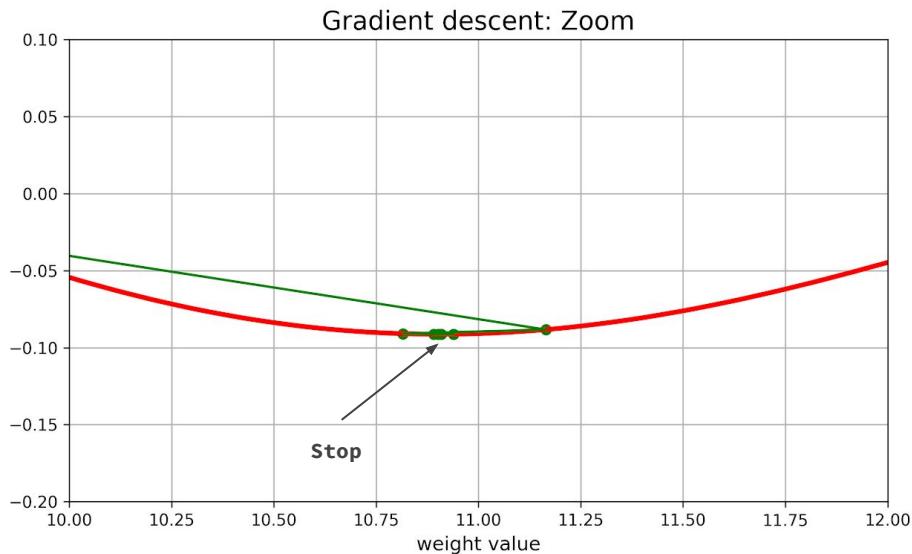
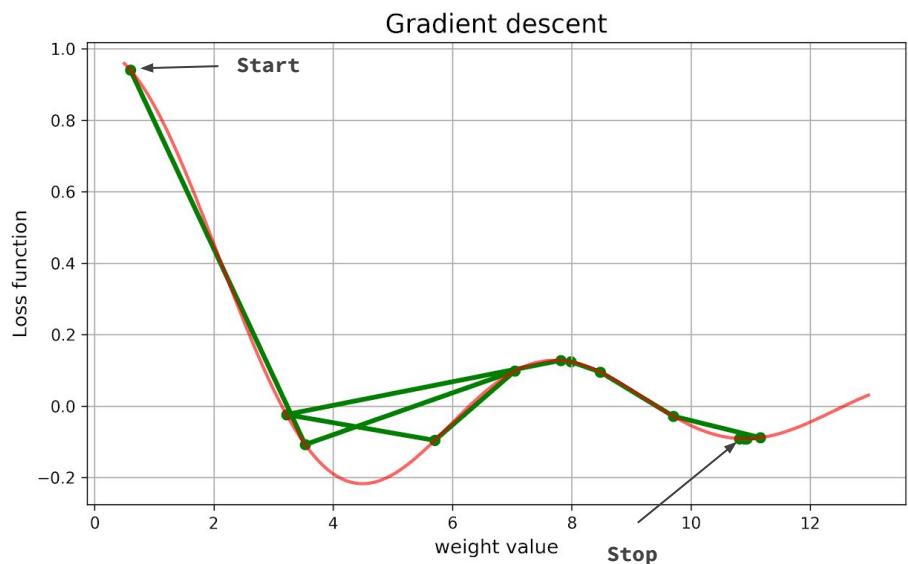
Optimization: Gradient descent



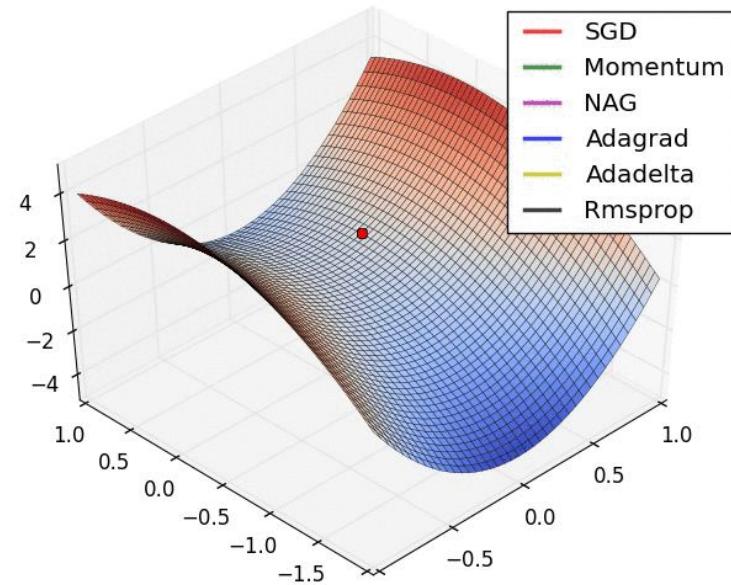
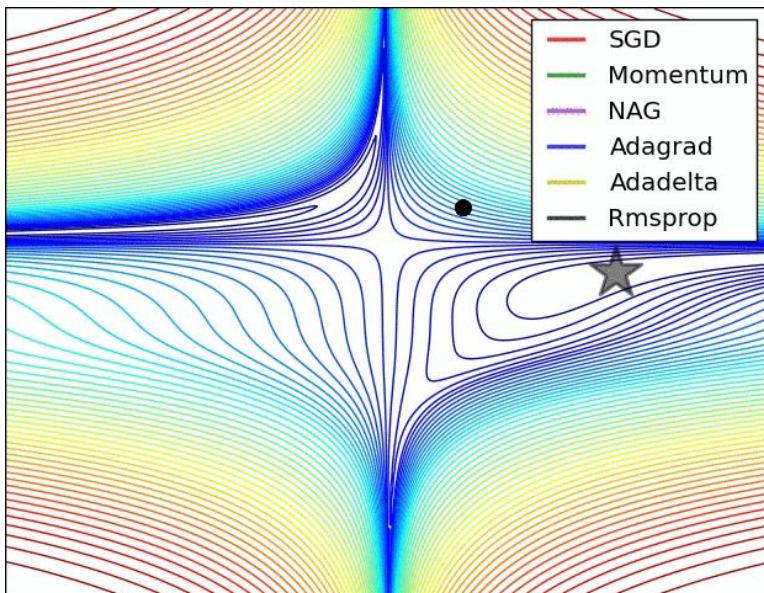
Optimization: Learning rate alto



Optimization: Problema del minimo locale



Optimization: Algoritmi oltre l'SGD



Nota: Immagini tratte da <http://cs231n.github.io/neural-networks-3/>

Backpropagation

- La Loss function in generale dipende da più variabili (weights e bias)

$$w(i+1) = w(i) - \eta \frac{dL}{dw} \rightarrow W(i+1) = W(i) - \eta \nabla L$$

$$W = (w_0, w_1, \dots, w_N)$$

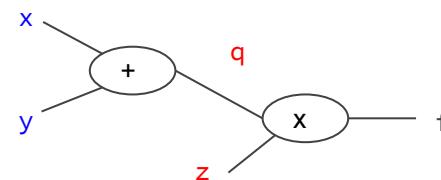
$$\nabla L = \left(\frac{\partial L}{\partial w_0}, \frac{\partial L}{\partial w_1}, \dots, \frac{\partial L}{\partial w_N} \right)$$

- **Chain rule:** $\frac{df(g(x))}{dx} = \frac{df}{dg} \frac{dg}{dx}$

$$\frac{df}{dq} = z; \quad \frac{df}{dz} = q$$

$$\frac{df}{dx} = \frac{df}{dq} \frac{dq}{dx}; \quad \frac{df}{dy} = \frac{df}{dq} \frac{dq}{dy}$$

$$f(x, y, z) = (x + y)z$$
$$f = qz; \quad q = x + y$$



Derivate locali

Derivate locali

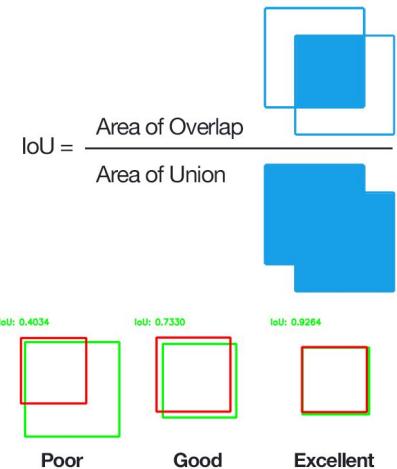
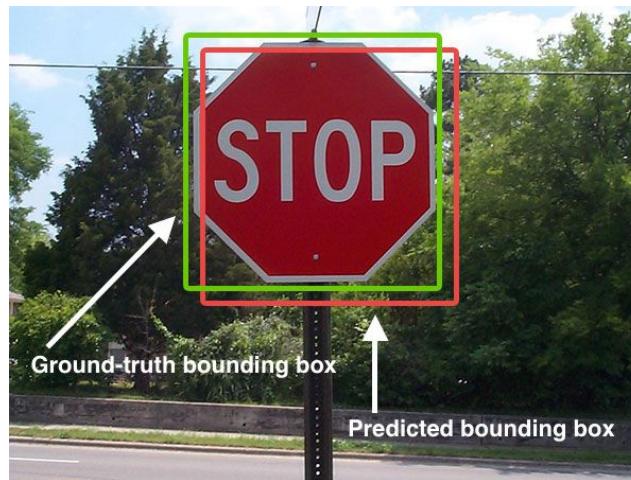
Misura della prestazione: Esempi

Accuracy



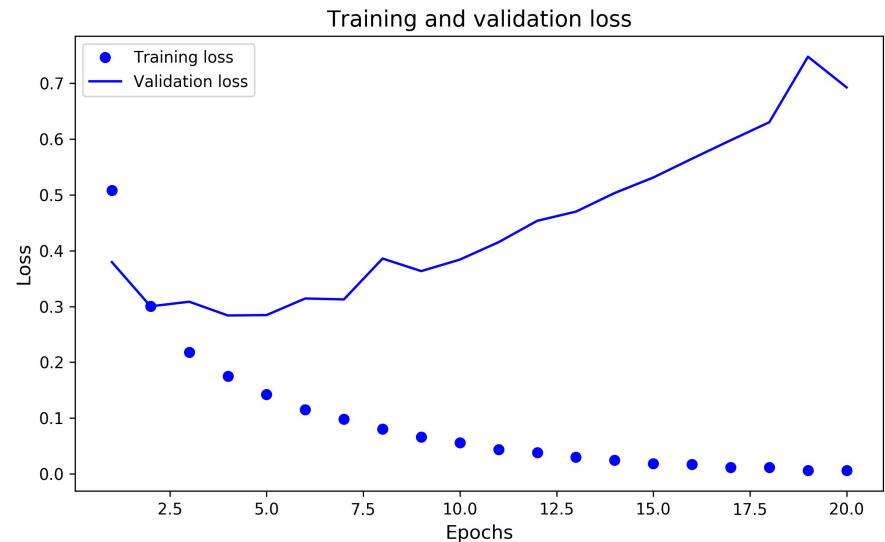
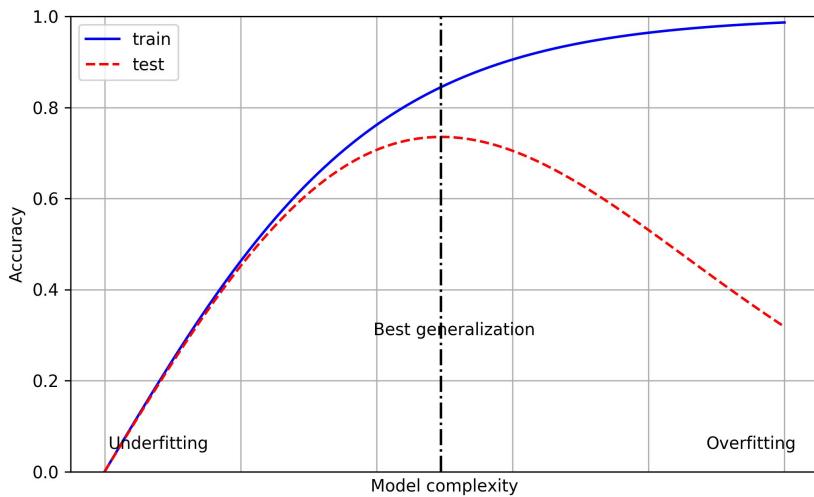
$$\frac{\text{Predizioni corrette}}{\text{Totale predizioni}} = \frac{TP+TN}{TP+TN+FP+FN}$$

Jaccard Index
(IoU)



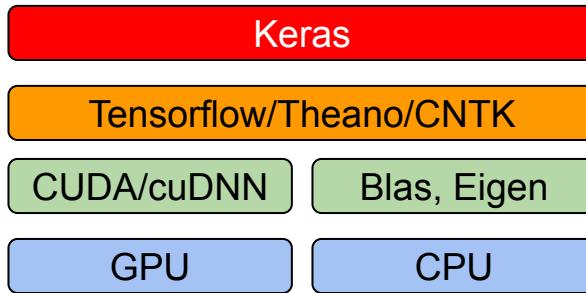
Overfitting: Validation set e test set

- Validation set e test set
- Tuning degli iper-parametri



Introduzione a Keras

Keras



- È una API Python di alto livello per poter sviluppare modelli con neural networks
- Può usare diversi tipi di Backend che si occupano dell'organizzazione dei calcoli
- Può sfruttare la potenza di calcolo offerta da diversi acceleratori (GPU, FPGA, TPU)

Keras: Tensore

- Nel contesto dei framework per DL/ML è un “contenitore” nel quale vengono strutturati i dati
- Può avere K dimensioni:
 - K=2: dati sotto forma di vettori → (samples, features)
 - K=3: timeseries → (samples, timesteps, features)
 - K=4: Immagini → (samples, channels, high, width)
 - K=5: Video → (samples, frames, channels, high, width)
- I layer di una ANN hanno come I/O un tensore

Keras: MNIST Example

- Database di cifre scritte a mano
- 60000 train, 10000 test
- Ogni dato di input è un'immagine a 8 bit da 28x28 pixels
- L'uscita desiderata è espressa con un intero (0-9)



Keras: MNIST, codice completo

```
from keras.datasets import mnist
from keras import models
from keras import layers
from keras.utils import to_categorical

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype('float32') / 255
test_images = test_images.reshape((10000, 28 * 28))
test_images = test_images.astype('float32') / 255
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

network = models.Sequential()
network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
network.add(layers.Dense(10, activation='softmax'))

network.compile(optimizer='rmsprop',
                 loss='categorical_crossentropy',
                 metrics=['accuracy'])

network.fit(train_images, train_labels, epochs=5, batch_size=128)

test_loss, test_acc = network.evaluate(test_images, test_labels)
print('test_acc:', test_acc)
```

Caricamento dei dati e preprocessing

Definizione rete

Training

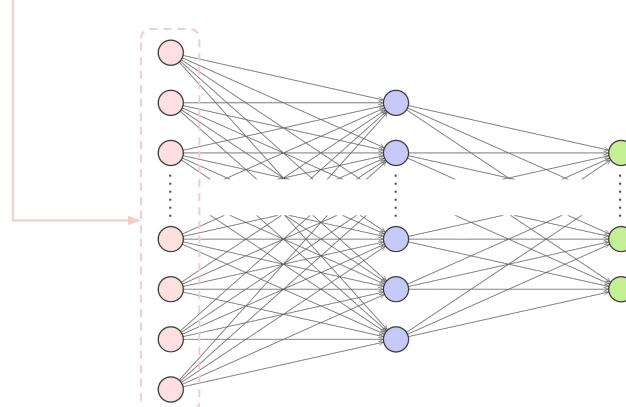
Testing

Keras: MNIST

Descrizione dell'architettura (Sequential model):

```
network = models.Sequential()  
network.add(layers.Dense(512, activation='relu', input_shape=(28*28,)))  
network.add(layers.Dense(10, activation='softmax'))
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
dense_2 (Dense)	(None, 10)	5130
<hr/>		
Total params: 407,050		
Trainable params: 407,050		
Non-trainable params: 0		



Keras: MNIST

Compilazione del modello e training (specifica loss, optimizer, metrics):

```
network.compile(optimizer='rmsprop',
                 loss='categorical_crossentropy',
                 metrics=['accuracy'])

network.fit(train_images, train_labels, epochs=5, batch_size=128)
```

Output fase di Training:

Epoch 1/5

60000/60000 [=====] - 4s 75us/step - loss: 0.2562 - acc: 0.9260

Epoch 2/5

60000/60000 [=====] - 5s 84us/step - loss: 0.1044 - acc: 0.9686

Epoch 3/5

60000/60000 [=====] - 5s 76us/step - loss: 0.0695 - acc: 0.9789

Epoch 4/5

60000/60000 [=====] - 5s 85us/step - loss: 0.0501 - acc: 0.9850

Epoch 5/5

60000/60000 [=====] - 4s 73us/step - loss: 0.0384 - acc: 0.9883

Keras: MNIST

Valutazione del modello (test set):

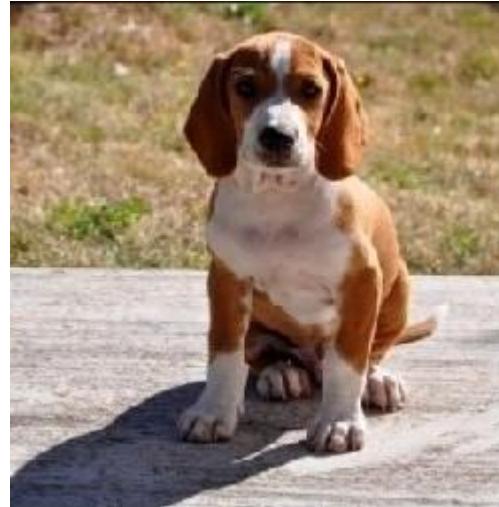
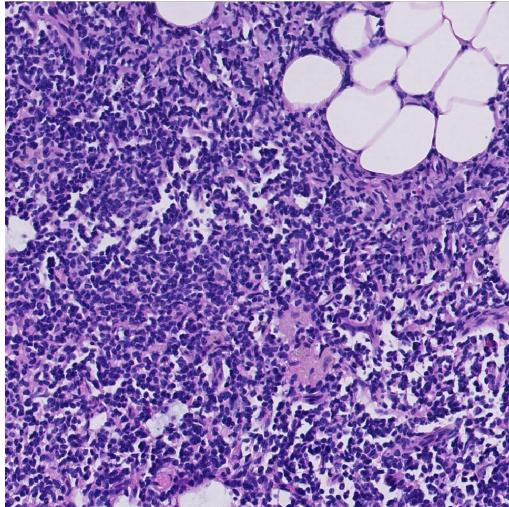
```
test_loss, test_acc = network.evaluate(test_images, test_labels)
print('test_acc:', test_acc)
```

10000/10000 [=====] - 0s 47us/step

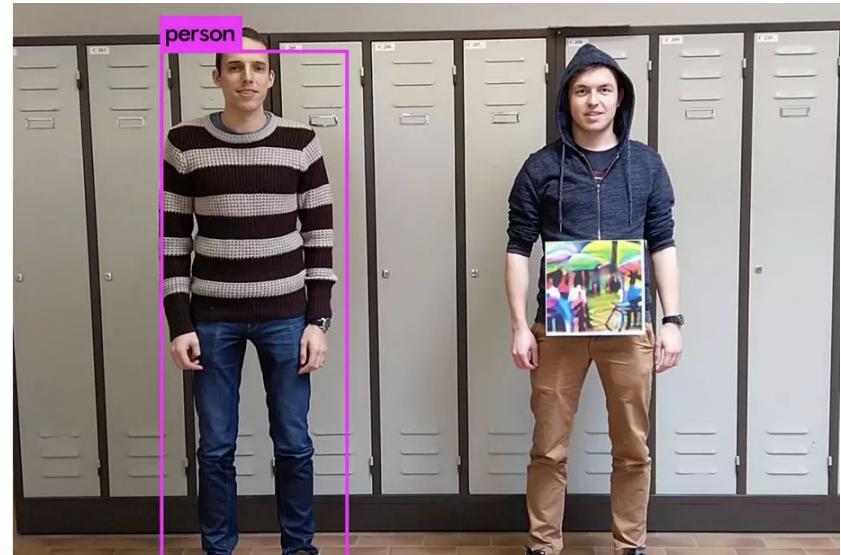
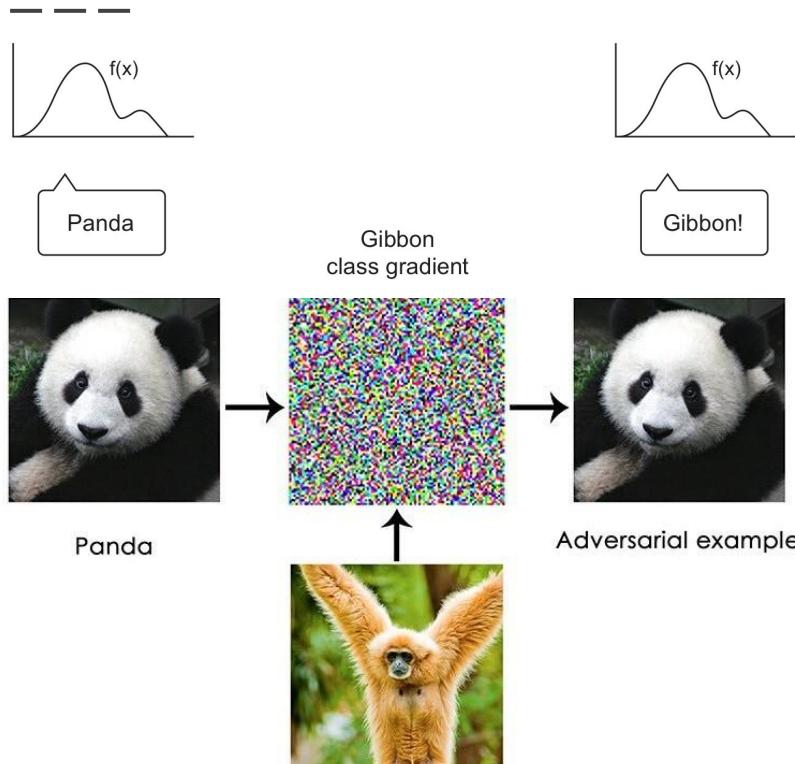
test_acc: 0.9774

Limiti del Deep Learning (1)

- La ANN riconosce solo quello che ha imparato



Limiti del Deep Learning (2): Adversarial images



Domande